

ZAYED ANSARI

AI Engineer

☎ +91-9356635531 ✉ zayed.ansari04@gmail.com 🔗 linkedin.com/in/zayed-ansari

Summary

Entry-level AI Engineer with experience building ML models and LLM-based applications using TensorFlow, PyTorch, Scikit-Learn, and related libraries. Comfortable working across data preprocessing, model training, evaluation, and deploying results via API(s) or simple interfaces for end-users. Seeking opportunities to design and ship practical ML solutions in real-world products.

Technical Skills

Languages & Frameworks: Python, Scikit-Learn, PyTorch, TensorFlow, LangChain, RAG, QLoRA

Tools & Libraries: Docker, MLflow, Git, NumPy, Flask, FastAPI

Databases: MongoDB, Redis, ChromaDB

Experience

Python Developer Intern

Jan 2026 - Apr 2026

Helo.ai by VivaConnect

Mumbai, Maharashtra

- Built a self-service MIS reporting POC where non-technical users generate reports via natural language prompts, integrating locally hosted Llama-3.2 (Ollama) with an agentic orchestration flow to output full reports from a single query.
- Automated schema change detection in the MIS reporting pipeline so that reports adapted to major or minor data changes in a few hours instead of days, significantly reducing manual rework for analysts.
- Developed a configuration-driven ETL pipeline in a Windmill OSS, three-node Python flow (extract, aggregate, save) operating across multiple MongoDB databases, fully API-triggered with zero hardcoded logic per data source.

Projects

Real-Time Fraud Detection Pipeline | *Kafka, PySpark, XGBoost* [Code]

Apr 2026

- Built an end-to-end streaming pipeline ingesting ATM and UPI transactions via Kafka through a Spark Structured Streaming job with an XGBoost UDF; scored each transaction in under 200 ms, achieving fraud precision of 0.85 and recall of 1.00 on 6.3M PaySim1 rows.
- Engineered a composite risk scoring engine layering ML score with behavioral signals, transaction velocity, new beneficiary detection, odd-hour flags, and balance drain ratio capping at 1.0 with a 0.60 threshold; Suspicious accounts were frozen in Redis with a 15-minute TTL.
- Implemented HMAC-SHA256 signed approve/reject email alerts with a fallback agent that auto-rejects on TTL expiry; built a live Flask dashboard auto-refreshing every 5 seconds with ML score, composite score, velocity, and approval state.

QLoRA Fine-Tuning on Llama 3.2 3B | *Unsloth, PyTorch, Hugging Face, Ollama*

Mar 2026

- Fine-tuned Llama 3.2 3B on FineTome-100k using QLoRA via Unsloth on a local GPU; applied gradient checkpointing, paged AdamW 8-bit optimizer, and 512-token sequence capping to fit within VRAM constraints.
- Achieved training loss reduction from 1.36 to 0.87 across the full run; exported the final checkpoint to GGUF format and loaded it into Ollama for local inference validation.

Multi-Document RAG with Source Citations | *LangChain, ChromaDB, Streamlit* [Code] [Post]

Jan 2026

- Built a production-ready RAG system for multi-document querying across PDFs with every response attributed to a specific source and page; benchmarked Ollama vs. Gemini API, selecting Gemini for measurably higher answer coherence.
- Implemented semantic search with ChromaDB embeddings and LangChain orchestration; Handled 27-page documents with context-aware retrieval, directly addressing hallucination in document Q&A.

Education

D Y Patil College of Engineering

Pune, Maharashtra

B.Tech Artificial Intelligence & Data Science | CGPA: 8.67

Nov 2022 - May 2026